



StackExchange Prediction and Analysis

Project for Advanced Data Mining Course

Tomasz Kawiak

Tomasz Makowski

Mateusz Mazur

December 20, 2025

Contents

Introduction	2
The Data	2
Data Characteristics and Preprocessing	2
Analysis of Answer Metrics and Sparsity	3
Tag Cardinality and Filtering	3
Question Score Distribution	4
Data gathering summary	5
The Methodologies	5
Embedding Generation	5
1. Global Document Embedding (Baseline)	5
2. Sequential Token Embedding	5
Embedding Analysis	5
Tag Dimensionality Reduction and Clustering	7
Optuna hyperparameter search	8
3. Recursive K-Means Clustering	10
4. Recursive Spherical K-Means Clustering	10
Reduction Analysis	11
The Experiments	12
Tag Prediction	12
XGBoost Baseline	12
Neural Network Architectures	12
Summary of Results	16
Score Prediction	17
Traditional ML Baseline	17
Deep Learning Regressors	17

Conclusions and Future Work	18
Key Successes and Insights	18
Summary of Best Results	18
Future Work	18
References	19

Introduction

This project investigates how deep learning and embedding methods can be applied in analysis and prediction of properties derived from StackExchange (StackOverflow) questions. Using the StackExchange API, we collected a dataset of 100,000 questions together with their metadata, textual context, and tag information. The analysis focuses on transforming the high-dimensional, sparse and highly imbalance data into representation suitable for learning.

A major technical challenge addressed was managing the extreme cardinality and high dimensionality of the data: specifically, compressing over 22,000 unique question tags into a meaningful, lower-dimensional space, and then leveraging 4096-dimensional embeddings for both multi-label classification (tag prediction) and high-variance regression (question score prediction). To accomplish this, we embedded all tags, titles and question bodies using the `qwen3-embedding:8b` model. For tags, we have applied a series of dimensionality reduction and clustering strategies. After evaluating UMAP (McInnes, Healy, and Melville 2020) + HDBSCAN (McInnes, Healy, and Astels 2017), Birch (Zhang, Ramakrishnan, and Livny 1997), and agglomerative approaches, we adopted **Recursive Spherical K-Means**, which produced 100 tag centroids that preserve coverage of all original tags.

The core objectives pursued throughout this project were:

1. **Data Exploration and Feature Engineering (EDA):** To quantify the sparsity and distribution characteristics of key features such as tag frequency, temporal metrics related to answer acceptance, and question scoring.
2. **Dimensionality Reduction:** To overcome computational limits and improve model tractability by intelligently reducing the space of 4096-dimensional embeddings and the semantic space of 22,753 unique tags.
3. **Tag Classification:** To design and evaluate neural network architectures capable of predicting question categories based on textual input from the title and body embeddings.
4. **Score Regression:** To assess the intrinsic predictability of question quality (measured by score) directly from the learned semantic embeddings.

The methodology relied heavily on the `qwen3-embedding:8b` model for vector representation and incorporates robust machine learning techniques such as **Recursive Spherical K-Means** for unsupervised clustering, and **Asymmetric Loss (ASL)** and **Cross-Attention** mechanisms for enhanced deep learning performance.

The Data

The analysis is based on a dataset comprising of **100,000 questions** extracted from the StackExchange platform (StackOverflow) using StackExchange API. The aim was to get to know the data characteristics and see what challenges may arise during modeling as well as what in particular can be predicted from the textual content.

Data Characteristics and Preprocessing

The raw dataset contained eleven distinct features covering textual content, metrics, and acceptance status:

#	Column	Type
0	<code>title</code>	<code>string</code>
1	<code>has_accepted_answer</code>	<code>bool</code>
2	<code>accepted_answer_score</code>	<code>float64</code>
3	<code>time_to_accepted_answer_hours</code>	<code>float64</code>
4	<code>question_score</code>	<code>int64</code>
5	<code>question_text</code>	<code>string</code>
6	<code>num_tags</code>	<code>int64</code>
7	<code>tags</code>	<code>string[]</code>

#	Column	Type
8	accepted_answer_id	float64
9	accepted_answer_length_chars	float64
10	accepted_answer_length_tokens	float64

Initial data hygiene involved dropping 8 duplicate questions identified by their ID.

Analysis of Answer Metrics and Sparsity

A crucial finding from the exploratory data analysis (EDA) was the significant sparsity in the answer-related features:

- A total of 39,938 questions had an accepted answer, while 60,054 did not.
- However, only **12,000** of these accepted answers had non-null values for temporal metrics.

Analysis of the `time_to_accepted_answer_hours` for this small subset revealed that the distribution, when log-transformed, exhibited a multimodal structure. This suggests that answers are accepted across different temporal regimes, possibly corresponding to simple versus complex problems, or different subject areas (fig. 1).




Figure 1: Gaussian Mixture Model Fit to Log Time to Accepted Answer

This result led to the conclusion that answer-based analysis was challenging due to the lack of sufficient data points with defined accepted answer characteristics. Scraping more data, just to analyze only a fraction of questions seemed infeasible, that's why we dropped this idea.

Tag Cardinality and Filtering

The raw dataset contained an excessively large vocabulary of **22,753 unique tags**. The top five most frequent initial tags were `python` (1528), `c#` (746), `javascript` (703), `c++` (689), and `java` (592). For simplicity sake, we initially focused on the 7,684 most frequent tags. The majority of questions in this subset had only one tag (38,547), although multi-label instances were present (3,015 questions had 2 tags; 464 had 3).

To create a manageable feature space for initial classification attempts, two approaches were considered:

1. **Frequency-Based Filtering:** Limiting the analysis to the most popular programming languages (e.g., Python, C#, Java). This approach resulted in a subset of 42,037 questions tagged with a set of 9 popular programming languages.
2. **Semantic Clustering:** Using high-dimensional embeddings and clustering algorithms to group semantically similar tags.

After a preliminary analysis, the second approach was chosen to retain semantic richness while reducing dimensionality and making the challenge more fun and interesting.




Figure 2: Tag frequencies for most frequent tags

Question Score Distribution

The raw scores (`question_score`) ranged widely from -20 up to 27,487. The score distribution was heavily skewed around zero (fig. 3). We considered grouping this continuous variable into five classes of uneven frequencies:

- **Bad:** $(-\infty, -1]$
- **Neutral:** 0
- **Good:** $(1, 3]$
- **Very Good:** $(3, 20]$
- **Excellent:** $(20, \infty)$

But ultimately, the decision was made to treat score prediction as a regression problem to preserve the granularity of the data.




Figure 3: Question Score Distribution

Data gathering summary

Finally, we retained the following columns for further analysis:

- `title` (textual content)
- `question_text` (textual content)
- `tags` (target for classification)
- `question_score` (target for regression)

The final version of the dataset contained **99 992 unique questions**.

The Methodologies

The path to generating robust predictive models required significant methodological investment, particularly in handling the immense size and complexity of the embedded text data.

Embedding Generation

To transform the textual data (titles, tags, and question bodies) into numerical feature representations, we utilized `qwen3-embedding:8b`, an open-source model capable of generating **4096-dimensional vectors** (Qwen Team 2025). Textual data (titles and question bodies) was transformed into **4096-dimensional vectors** using the open-source embedding model `qwen3-embedding:8b`.

We implemented two distinct embedding strategies:

1. Global Document Embedding (Baseline)

- This approach served as the baseline due to its simplicity and relative computational speed.
- Embed each question body, title and tag individually into an embedding vector.
- This method assumes that the semantic context of the entire document can be effectively compressed into a single 4096-dimensional vector (using `float64` precision) without significant information loss.
- Computation required approximately 6 hours using the `ollama` library (Ollama 2025). The resulting dataset occupied 3.3 GB of storage.

2. Sequential Token Embedding

- To address potential information loss in the baseline approach, we hypothesized that a single vector might fail to capture complex dependencies in longer texts.
- Instead of pooling the text into one vector, we maintained a sequence of embeddings to preserve token-level knowledge. We defined a fixed sequence length of 4 tokens for the title and 32 tokens for the body, resulting in a distinct embedding vector for each token.
- Implementation changes:
 - Due to the limitations of `ollama` in the terms of appropriate tokenizer for our model, as well as inefficient resource management during embedding, we've switched to *Hugging Face transformers* library.
 - Because of the exponential increase in data size and compute time, we reduced the floating-point precision from `float64` → `float32`.
- These optimizations reduced the estimated compute time from 40 hours to approximately 13 hours.
- The resulting embeddings required 27 GB of space. Due to memory constraints preventing the dataset from being loaded entirely into RAM, we utilized the Hierarchical Data Format (HDF5) for efficient storage and access (The HDF Group 1997-2025).

Embedding Analysis

After the embedding phase, we wanted to verify the validity of the topology and density of the resulting embedding space. Specifically, we wanted to verify that the embeddings captured sufficient semantic overlap between questions to facilitate meaningful clustering.

To measure this, we analyzed the *Nearest-Neighbour Cosine Similarity*. Using a subset of the question embeddings, we performed the following steps:

- Normalized the embeddings (the resulting embeddings from `ollama` theoretically should be normalized, but it's better to normalize it anyway, since it does not affect the data).
- We utilized the `NearestNeighbors` algorithm (from `scikit-learn`) to locate the closest non-identical neighbor ($k = 1$) for each data point.
- We calculated the cosine similarity for these pairs, defined as $1 - \text{cosine_distance}$.




Figure 4: Nearest-Neighbour cosine similarity distribution

As we can see in Figure fig. 4, the distribution of nearest-neighbor similarities is approximately Gaussian and centered around 0.65.

- The lack of data points near 0.0 indicates that very few questions are “isolated” in the vector space; almost every question has a semantically related counterpart.
- The unimodal distribution suggests a well-structured manifold where local neighborhoods are consistent. This confirms that the embedding model (`qwen3-embedding`) successfully mapped semantically similar questions to adjacent regions in the high-dimensional space, providing a strong foundation for the subsequent clustering phases.

Tag Dimensionality Reduction and Clustering

The tag space presents two major challenges: high dimensionality (4096-dimensional embeddings) and high cardinality (22,753 unique tags). These properties make traditional clustering methods difficult to apply directly. Because we aimed to preserve as much semantic structure as possible, we investigated several dimensionality reduction and clustering strategies.

1. HDBSCAN

HDBSCAN (Hierarchical Density-Based Spatial Clustering of Applications with Noise) (McInnes, Healy, and Astels 2017) was initially considered, due to its ability to identify clusters of varying density and to naturally model noise. However, the method proved infeasible at our scale. During graph construction, the algorithm attempted to allocate a dense distance matrix, resulting in more than 64 GB of RAM usage:

```
MemoryError: Unable to allocate 74.5 GiB for an array with shape (99992, 99992) and data type
→ float64
```

Even incorporating Birch pre-clustering (Zhang, Ramakrishnan, and Livny 1997), which is often recommended to reduce memory footprint, did not sufficiently mitigate these requirements.

2. UMAP

Given its popularity for high-dimensional manifold learning and its strong community reports, particularly its successful use in document embedding tasks such as the [20 Newsgroups dataset](#), we next explored UMAP (Uniform Manifold Approximation and Projection) (McInnes, Healy, and Melville 2020) as a potential solution. Encouraged by these findings, following community best practices, we tested two approaches:

- Single stage with ‘pure’ UMAP
- Two-stage with PCA → UMAP

The rationale for the second approach comes from the UMAP’s reliance on pairwise distances when constructing its topological graph. In very high-dimensional spaces, distance measures tend to concentrate, causing points to appear nearly equidistant. Applying PCA first captures the dominant variance structure and reduces sparsity effects before manifold learning.

For the single stage approach, we applied UMAP directly to the 4096-dimensional vectors using the cosine metric (appropriate given that the underlying embeddings were already normalized thanks to [ollama](#)). Reducing directly to two dimensions produced the projection shown in fig. 5, which exhibited significant overlap and very poor global structure.




Figure 5: Naïive UMAP reducing tag embeddings to 2 components (colors are just for aesthetic reasons).

After brief hyperparameter search using optuna, we've decided to scrap the single-stage approach, in favor of the two-stage one.

As the first stage, PCA was picked as the best candidate, due to it's variance maximization, relative simplicity, and little computation overhead.

Optuna hyperparameter search

To improve UMAP's performance, we conducted a hyperparameter search using Optuna (Akiba et al. 2019). For optimization, a suitable objective function was required. We first experimented with the Calinski–Harabasz index (Caliński and Harabasz 1974), but due to its tendency to assign inflated scores to degenerate solutions (e.g., one or two dense clusters), we replaced it with a weighted combination of:

- Silhouette Score
- Calinski-Harabasz score
- Custom penalty for generating clusters with very few points.

We tested both a flat optimization procedure and a nested one. In the nested variant, the outer loop optimized PCA hyperparameters, and the inner loop optimized UMAP parameters conditioned on the PCA output.

Neither approach produced coherent clustering. The resulting label distributions (fig. 6) were dominated by either a large noise cluster or numerous trivial micro-clusters. This instability pointed to deeper issues in the reduction pipeline:

- The evaluation metrics themselves were insufficiently sensitive to the structure we hoped to recover.
- Available compute limited the ability to explore more.
- Most importantly, the underlying assumptions of UMAP were violated:
 - UMAP assumes the data is sampled from a uniformly distributed Riemannian manifold, but dense embedding spaces often exhibit highly non-uniform local geometry.

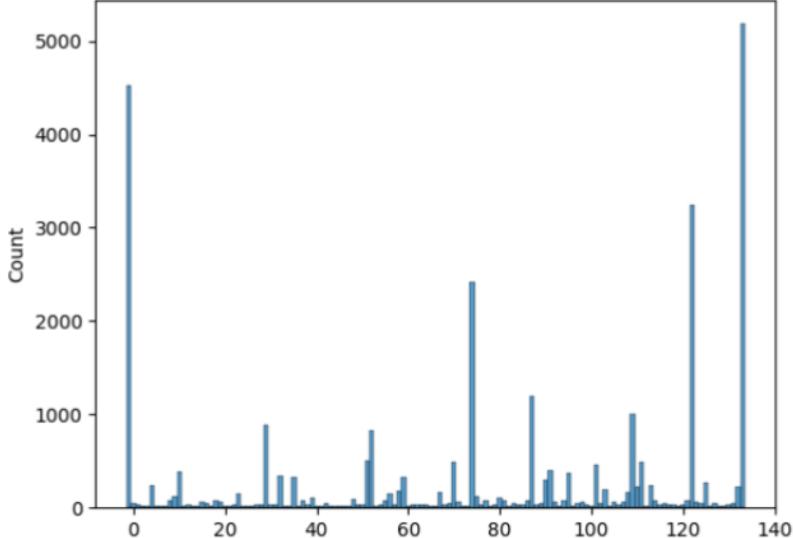


Figure 6: Histogram of labels resulting from UMAP reduction coupled with HDBSCAN. Each y value corresponds to a single cluster.

Variance Loss in PCA

A more fundamental limitation emerged when examining the PCA stage itself. To measure how much information a PCA stage would discard, we computed the explained variance on a representative sample of 10,000 embeddings. The results revealed that the embedding space is extremely flat: the first principal component alone explains a negligible fraction of the variance, and adding more components yields slow, sublinear improvement. Even with the top 50 PCs, only 38% of the total variance is recovered. This behavior is illustrated in fig. 7.

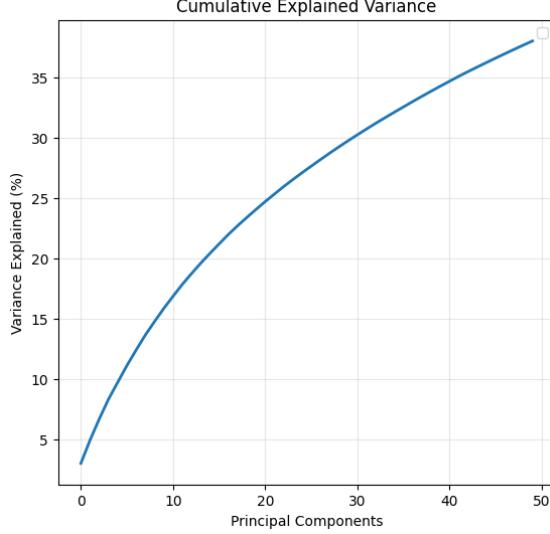


Figure 7: Cumulative variance curve of PCA.

This structure is also visible when comparing the distribution of a raw embedding dimension to the distribution of the first principal component. As shown in fig. 8, a random raw dimension is tightly concentrated around zero, indicating very low variance. In contrast, PC_1 has a much broader distribution, reflecting the fact that most of the meaningful variance in the data collapses.

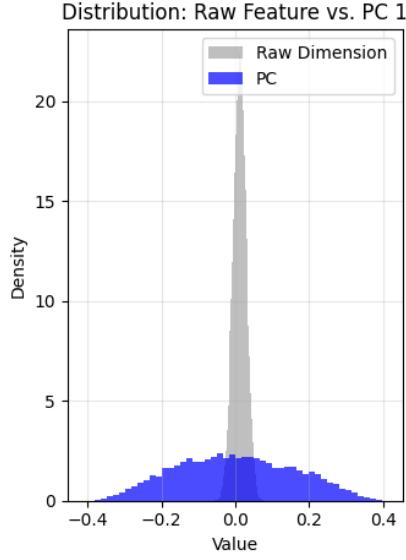


Figure 8: Cumulative variance curve of PCA.

To reiterate, metric limitations, computational constraints, and mismatches between the UMAP’s assumptions and the geometry of dense embedding space, combined with the extensive optuna hyperparameter search, for which the best candidate identified by the optimization process, the PCA alone discarded approximately 62% of the variance present, motivated us to search for different approach to dimensionality reduction.

3. Recursive K-Means Clustering

Our initial approach was inspired by methodologies for hierarchical data organization, such as those described in engineering blogs by [Spotify](#) dealing with large-scale recommendation systems. We designed a Recursive K-Means Clustering algorithm to structure the tags into a navigable tree.

The core concept involves recursively partitioning the tag embedding space. Starting with the root node containing all tags, we apply K-Means to split the data into k clusters. This process is repeated for each resulting cluster until a maximum depth is reached or a cluster becomes too small to split further. For this algorithm, we've relied on Euclidean distance.

The big difference from the previous algorithms is that we do not reduce the data itself, but rather **find the best representation for the predefined amount of classes**.

Tag	closest_dist	cluster_mean	tag_embedding	cluster_size	top_5_popular_tags
facebook	0.089584	[0.004879954, ...]	[0.002688237, ...]	94	[instagram-api, instagram, ...]
integer	0.040206	[0.027635492, ...]	[0.03159611, ...]	370	[intervals, differential-equations, ...]
storage	0.105308	[0.008282, ...]	[0.020607475, ...]	134	[dicomweb, remotestorage, ...]
tf-idf	0.070302	[0.027263727, ...]	[0.024365697, ...]	243	[wpml, locale, ...]
qt	0.074746	[0.01614129, ...]	[0.016203284, ...]	106	[qt, qml, pyqt, qasync, qt6]

4. Recursive Spherical K-Means Clustering

While effective for spatial data, Euclidean distance is often suboptimal for high-dimensional semantic embeddings, where the direction of the vector typically encodes more semantic meaning than its magnitude. That's why, to better capture the semantic relationships between tags, we evolved our approach to Recursive Spherical K-Means Clustering (Hornik et al. 2012). This variation partitions the data by minimizing the cosine dissimilarity rather than the Euclidean distance, effectively clustering data points on the surface of a hypersphere (hence, the spherical in Spherical K-Means).

Implementation Details

After investigating the tag frequency in our dataset, we've decided to impute all tags, which occurred less frequent than 100 times. This resulted in 411 unique tags and in 90205 questions with tag. This operation did not change the overall characteristic of the dataset as we can see in the fig. 9




Figure 9: Distribution of number of tags per question

With the orphan questions (the ones without tags), we assigned them to the nearest embedding based on the cosine similarity of their tag.

For example:

```
closest_centroid_tag(tags,centroid_tags, 'ntp')
('datetime', np.float64(0.8098148848579392))
```

We implemented a custom SphericalKMeans estimator compatible with the *scikit-learn* API. The algorithm enforces unit-norm constraints on both the input data and the cluster centroids during optimization. The update steps were optimized following the algorithms described by (Schubert, Lang, and Feher 2021).

Our implementation include:

- Unit Normalization
- Cosine Similarity maximization:
 - The core objective is to maximize the summation of cosine similarities between samples and their assigned centroids.
- The algorithm returns a tree structure. The root node contains all tag embeddings; the next level (depth 1) contains the target centroids (100 in our case); subsequent levels contain increasingly specific tag embeddings.
 - For all target tags, we calculated the mean and identified the ‘most representative tag’ (the tag closest to the cluster mean).

From this point forward, we refer to the clusters at depth 1 of the recursive tree as the clusters for simplicity sake.

Reduction Analysis

The resulting clusters demonstrated strong semantic coherence. For example, for the outlier centroid with most directly connected tags, the algorithm effectively grouped database-related technologies:

- Cluster Centroid: `database`
- Members: `database, mongodb, sql, elasticsearch, firebase, postgresql, sqlite, mysql, jdbc, supabase, sql-update, redis, hibernate, sqlalchemy, snowflake-cloud-data-platform`

Similarly, we observed distinct clusters for front-end frameworks (grouping angular, vue.js, react), mobile development (android, flutter, ios), and scientific computing (pandas, numpy, python).

Another impressive result of this reduction is its ability to classify question embeddings solely based on cosine similarity. As an example, we evaluated how well the model classifies the `.net` tag. We computed the cosine similarity between a given question and each cluster, selecting the two closest clusters. We then computed a confusion matrix based on the presence of `.net` in the ground-truth question tags, which yielded the results in fig. 10.




Figure 10: Confusion Matrix of top 2 results for `.net`

Based on this functionality, we've made a basic model `DualEncoderMatcher`, which classifies tags based on cosine similarity mentioned above.

As an experiment, we took random question and checked 5 most similar clusters to it:

- Question:

```
'I am new to pandas library in python. When I loaded a file and was printing the output of df.info  
→ into console, the data ...'
```

- Result:

```
Top 5 predictions: [('python', np.float64(0.38295367797329094)), ('word-table',  
→ np.float64(0.3696944707340029)), ('ironpdf', np.float64(0.3647400168217892)), ('django',  
→ np.float64(0.3297345619706671)), ('apache-spark', np.float64(0.3283243442282066))]
```

These experiments demonstrated the semantic power held by these clusters and prompted us to move towards prediction.

The Experiments

Tag Prediction

The primary classification goal was to predict the appropriate semantic tag centroid(s) for a given question using its 4096-dimensional body embeddings.

XGBoost Baseline

As a baseline, an XGBoost (Chen and Guestrin 2016) model was trained on question body embeddings. To simplify the initial evaluation, the inherently multi-label task was reduced to a multiclass classification problem. The target for each question was mapped to a single label via a majority vote mechanism, selecting the centroid (or tag group) containing the highest frequency of the question's original tags.

This XGBoost model trained for 599 minutes. It achieved a **Weighted F1 Score of 0.6882** (and accuracy of 0.6928). The limitation of this approach was its inability to effectively model the complex, non-linear semantic interactions embedded in the 4096-dimensional vectors, forcing it to treat the dimensions largely independently. Nevertheless, this result provided a solid benchmark for subsequent deep learning models.

Neural Network Architectures

To improve baseline performance, several neural network architectures were explored to better capture the semantic relationships in the embeddings and to natively handle the multi-label nature of the task.

Simple MLP Baseline

We implemented a baseline using a simple Multi-Layer Perceptron (MLP) trained on question body embeddings. Given the multi-label nature of the problem, where a single question may contain both popular tags (e.g., `python`) and rare tags (e.g., `darts`), we had to employ a loss function which takes it into consideration. Initially, we employed `BCEWithLogitsLoss`, which applies a binary cross-entropy loss independently to each class.

While `BCEWithLogitsLoss` is standard for multi-label tasks, it treats all negative labels equally. In a sparse setting where most tags are negative for any given sample, the easy negatives can overwhelm the training signal. Despite these limitations, the MLP baseline achieved a weighted F1 score of approximately 0.6790. This reasonable performance confirmed the viability of the embedding approach but highlighted the need for richer architectures to capture the nuanced relationship between a question's title and its body.

Dual-Stream Fusion Network (DSF)

To better leverage the distinct semantic information contained in titles and bodies, we developed a custom Dual-Stream Fusion (DSF) network, adapting the architecture proposed by Yang et al. for multi-modal fake news detection (Yang et al. 2024).

Architecture The model processes the Title and Body embeddings through two separate streams. These streams are then integrated using a Multi-Head Self-Attention (MHSA) mechanism. The core idea is that the attention mechanism can dynamically weigh the importance of the title versus the body depending on the context. For instance, a short, distinct title might carry more weight than a long, vague body text.

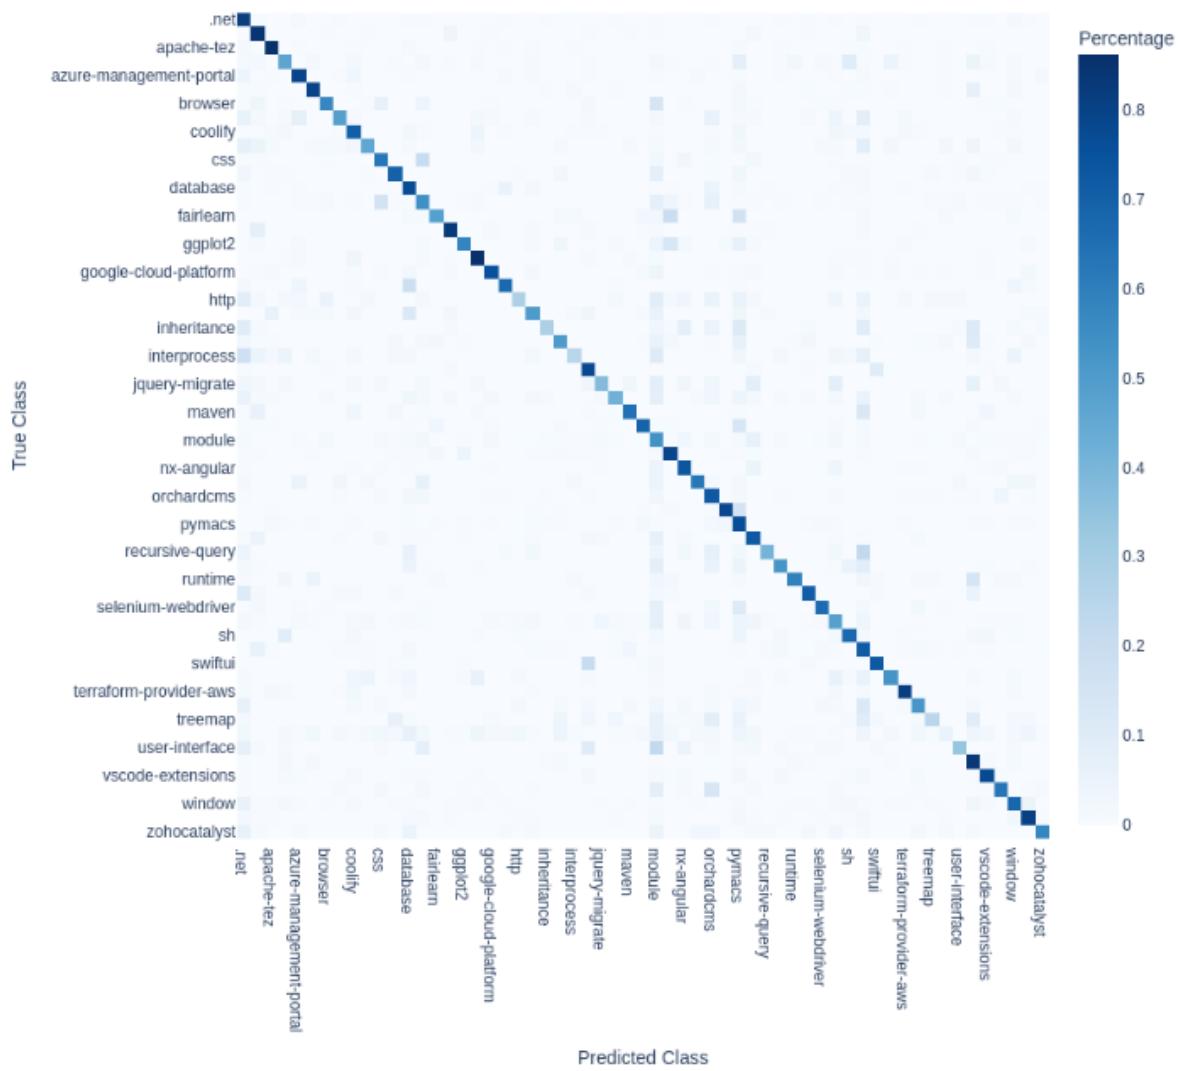


Figure 11: Confusion Matrix of XGBoost Classifier

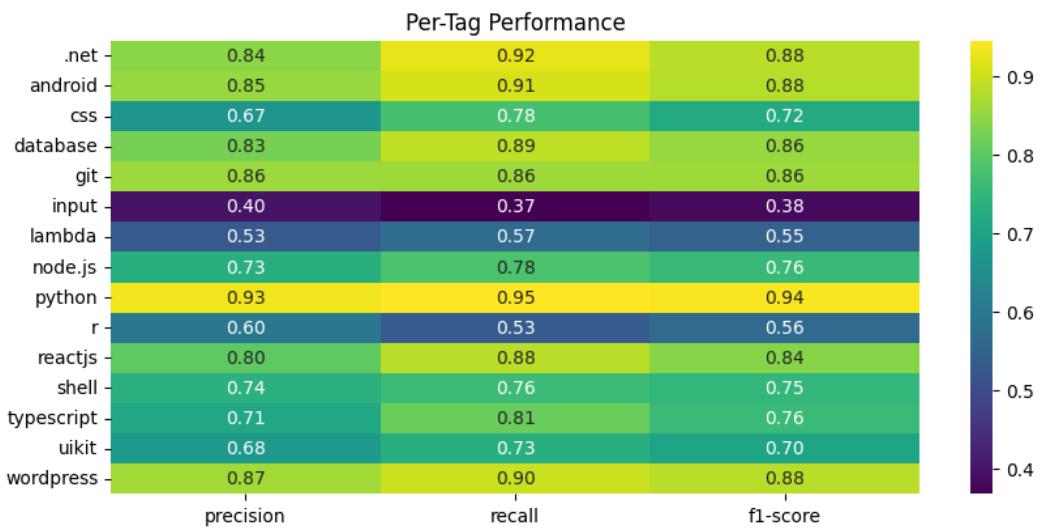


Figure 12: Validation set performance for MLP baseline model

Regularization and Overfitting Initial training runs revealed significant overfitting. The model’s validation performance rapidly deteriorated below the MLP baseline after only a few epochs. This demonstrated the “double-edged sword” of attention mechanisms: while they offer high expressivity, they allow the model to easily memorize the relatively small dataset of $\sim 100k$ samples. Standard regularization techniques—such as increased dropout, weight decay, and aggressive learning rate schedulers—yielded limited success.

Notable improvement came from implementing a different loss function—**Assymmetric Loss** (ASL) (Ben-Baruch et al. 2021). In multi-label classification with many classes, the “negative” samples (tags not present) vastly outnumber the “positive” ones. Standard Cross Entropy allows these easy negatives to dominate the gradient, washing out the signal from the rare positives. ASL addresses this by dynamically down-weighting easy negatives. It introduces two focusing parameters, γ_+ and γ_- . By setting $\gamma_- > \gamma_+$, ASL aggressively suppresses the loss contribution from negative samples the model is already confident about. This forces the optimization process to focus on “hard” negatives (confusing tags) and positive samples, effectively handling the extreme class imbalance without manual re-weighting.

We trained the optimized DSF model with **AdamW** optimizer, **ReduceLROnPlateau** scheduler, and ASL loss with $\gamma_- = 4, \gamma_+ = 1$ for 50 epochs, which resulted in F1 score (micro) on validation set of 0.713, and weighted of 0.711.



Figure 13: Snippet of Per-Tag performance of DSF-MHSA model on the validation set. Each column reflects precision, recall f1-score respectively.

As seen in fig. 13, the use of a high γ_- in ASL successfully boosted recall for many classes. However, this comes with a trade-off: simply increasing γ_- indefinitely leads to an overemphasis on recall at the expense of precision. This behavior indicates that while advanced loss functions can mitigate data imbalance, further performance gains likely require increasing the dataset size rather than just architectural tuning.

2. DSF with Cross-Attention Fusion (Optimal Model):

- To further combat overfitting and improve feature interaction, the MHSA was replaced with **Cross-Attention**, allowing the typically concise **Title embedding to “query” the verbose Body embedding** to extract relevant features.
- Additional regularization, **Manifold Mixup**, was applied to the embeddings during training to encourage smoother decision boundaries in the latent space.
- Trained over 100 epochs (~ 60 minutes) using OneCycleLR scheduling, this model achieved the best outcome: **F1 Micro 0.7253** and **F1 Weighted 0.7196**.

Seq2Seq Model

Instead of thinking about our task as a classification task, we could rephrase it as a Sequence to Sequence problem, where we want to model one sequence (question embedding) into another (tag(s) embedding(s)). This approach has the biggest potential out of all mentioned earlier, since we can leverage the unprecedented rise of LLMs, by taking a pre-trained model (in our case `t5-small`) and fine-tune it for our task.



Figure 14: Per tag performance of DSF with Cross-Attention Fusion

The biggest advantage is also its biggest disadvantage—all modern models require prohibitively large amount of compute, which forced us to pick a relatively small (60M) model. We picked `t5-small` because it was trained primarily on summarization and translation task, which is exactly what we want this model to do.

After 5 hours of training, we've achieved F1 micro score of 0.6676 and F1 macro of 0.625. While this may not sound as impressive as previous models, we have to keep in mind, the size of this model, and the transformers innate need for huge amount of data.




Figure 15: Training Loss and Validation F1 over Steps

Another advantage of this model is the ease of use in terms of user readable format. Using the `transformers` we can easily create quick predicting function for any given question out of data:

```
t = "How do I reverse a list?"
b = "I have a list [1, 2, 3] and I want [3, 2, 1]. slicing doesn't work for me. I'm thinking of
    ↪ using a pandas library"
print(predict_custom_question(t, b))
dataframe, python, sorting
```

What's also powerful, is that the model itself infers how many tags are needed for each question.

Summary of Results

As shown in [tbl. 3](#), the DSF with Cross-Attention Fusion outperformed all other models, demonstrating the effectiveness of attention-based fusion and advanced loss functions in multi-label tag prediction tasks.

Table 3: Summary of Tag Prediction Results.

Model	F1 Score (Weighted)
XGBoost (Multiclass Approximation)	0.6882
Baseline MLP	0.6790
DSF with MHSA Fusion	0.7110
DSF with Cross-Attention Fusion	0.7196

The final model mainly has trouble with tags like “import,” “installation,” and “validation,” which probably suggests that it is difficult to distinguish between common technical noise and particular topical intent. The Asymmetric Loss (ASL) may have over-suppressed terms like “import” as “easy negatives” because they appear as boilerplate in nearly every code snippet. Furthermore, the model’s inability to handle specialized tags like “asp.net-web-api” indicates that the Cross-Attention mechanism may occasionally lack a detailed enough Title “Query” to extract particular nuances from the verbose Body text.

Score Prediction

Our secondary regression task focused on predicting the question score, which reflects community engagement and perceived quality. The goal was to predict the raw integer question score using only the embedded textual content (regression task). This task was inherently challenging due to the high variance and sparse nature of scores (mean score of 23.55, but max score of 27,487, with most scores clustered near zero).

Traditional ML Baseline

To establish a performance floor, we conducted an initial exploration using traditional NLP techniques, combining TF-IDF feature extraction with Truncated SVD (Latent Semantic Analysis) for dimensionality reduction. We benchmarked several configurations, including:

- Linear Models: TF-IDF (with variations in n-grams and stemming) paired with Ridge Regression.
- Ensemble Methods: Gradient Boosting (GBR) utilizing TF-IDF, SVD, and engineered “extra features.”
- Baseline Comparisons: Standard Bag-of-Words (BoW) and a simple mean-prediction baseline.

The results were underwhelming:

- The mean baseline achieved an R^2 of -0.000196 (Test RMSE 149.97).
- The best traditional model, **TF-IDF + SVD + Ridge**, managed a Test R^2 of **0.0074** (Test RMSE 275.05) on the raw score target.

This near-zero R2 score underscores the inherent difficulty of the task: traditional frequency-based features are insufficient for capturing the complex, non-linear relationships that drive community engagement (scores) on Stack Overflow. This served as a strong justification for moving toward the deep learning approaches detailed below.

Deep Learning Regressors

To evaluate the predictive power of neural architectures on numerical outcomes, we adapted the single-stream MLPs and the dual-stream DSF variants for regression by utilizing a Mean Squared Error (MSE) loss function. In addition to the semantic embeddings, we integrated a normalized, clipped tag count feature to provide the model with a proxy for question complexity.

As shown in [tbl. 4](#), all deep learning configurations achieved a substantial performance leap over the traditional baselines ($R^2 \approx 0.000$).

Table 4: Score Regression Performance on Test Set (using Embeddings).

Model	Input	Test RMSE	Test R^2
Mean Baseline	Constant	149.97	0.000
Title MLP	Title Embedding Only	138.73	0.144
Body MLP	Body Embedding Only	137.57	0.158
DSF-CrossAttention Regressor	Title + Body Embeddings	135.27	0.186
DSF-MHSA Regressor	Title + Body Embeddings	134.24	0.199

The DSF-MHSA Regressor emerged as the top performer, explaining approximately 19.9% of the score variance. Interestingly, in the regression context, the global focus of Multi-Head Self-Attention (MHSA) slightly outperformed the more targeted Cross-Attention mechanism.

Despite the improvement, diagnostic analysis revealed a common challenge in social data regression: while the model accurately predicts the vast majority of low-scoring posts, it struggles to capture the “viral” outliers or high-score peaks ([fig. 16](#)). This suggests that high scores may be driven by external temporal factors or community dynamics not fully captured within the text embeddings alone.




Figure 16: Predicted vs Actual Scores (DSF-MHSA, Filtered Range)

Conclusions and Future Work

The overall project successfully navigated several data processing and modeling complexities inherent in large-scale textual data.

Key Successes and Insights

- **Robust Tag Reduction:** The greatest technical victory in the data preparation phase was the establishment of the **Recursive Spherical K-Means** pipeline, effectively reducing the 22,753 original tags to 100 semantically coherent centroids. This step was vital for making multi-label classification computationally feasible after the failure of unsupervised methods like UMAP/HDBSCAN optimization.
- **Optimal Classifier Architecture:** The use of **Dual-Stream Fusion Networks** demonstrated a clear advantage over both traditional ML (XGBoost) and simple MLPs for multi-label classification. The best results were achieved by the **DSF with Cross-Attention Fusion (F1 Weighted 0.7196)**, incorporating both **Asymmetric Loss** to manage label imbalance and **Manifold Mixup** to improve generalization.
- **Score Prediction Difficulty:** The consistently low R^2 values across all regression experiments confirm that predicting Stack Overflow scores from purely semantic content embeddings is inherently difficult due to the social/external factors influencing a question’s eventual popularity (score).

Summary of Best Results

Task	Best Model	Key Metric	Result
Tag Prediction (Classification)	DSF Cross-Attention	F1 Weighted	0.7196
Score Prediction (Regression)	DSF-MHSA Regressor	Test R^2	0.199

Future Work

Future research should focus on mitigating the score prediction problem by incorporating features that capture non-textual quality signals, such as user reputation or time-of-day posting biases, which were outside the scope of this content-based embedding analysis. For classification, exploring even deeper transformer architectures for sequence-aware embedding generation could further enhance performance beyond the current pooled embedding methods.

References

- Akiba, Takuya, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. 2019. “Optuna: A Next-Generation Hyperparameter Optimization Framework.” In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*.
- Ben-Baruch, Emanuel, Tal Ridnik, Nadav Zamir, Asaf Noy, Itamar Friedman, Matan Protter, and Lihi Zelnik-Manor. 2021. “Asymmetric Loss for Multi-Label Classification.” <https://arxiv.org/abs/2009.14119>.
- Caliński, T., and J Harabasz. 1974. “A Dendrite Method for Cluster Analysis.” *Communications in Statistics* 3 (1): 1–27. <https://doi.org/10.1080/03610927408827101>.
- Chen, Tianqi, and Carlos Guestrin. 2016. “XGBoost: A Scalable Tree Boosting System.” In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–94. KDD ’16. ACM. <https://doi.org/10.1145/2939672.2939785>.
- Hornik, Kurt, Ingo Feinerer, Martin Kober, and Christian Buchta. 2012. “Spherical k-Means Clustering.” *Journal of Statistical Software* 50 (10): 1–22. <https://doi.org/10.18637/jss.v050.i10>.
- McInnes, Leland, John Healy, and Steve Astels. 2017. “Hdbscan: Hierarchical Density Based Clustering.” *The Journal of Open Source Software* 2 (March). <https://doi.org/10.21105/joss.00205>.
- McInnes, Leland, John Healy, and James Melville. 2020. “UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction.” <https://arxiv.org/abs/1802.03426>.
- Ollama. 2025. “Ollama: Get up and Running with Llama 3.2, Mistral, Gemma 2, and Other Large Language Models.” <https://github.com/ollama/ollama>.
- Qwen Team. 2025. “Qwen3-Embedding-8B: A Versatile Text Embedding Model.” <https://huggingface.co/Qwen/Qwen3-Embedding-8B>; Hugging Face.
- Schubert, Erich, Andreas Lang, and Gloria Feher. 2021. “Accelerating Spherical k-Means.” In *Similarity Search and Applications*, 217–31. Springer International Publishing. https://doi.org/10.1007/978-3-030-89657-7_17.
- The HDF Group. 1997-2025. “Hierarchical Data Format 5 (HDF5).” <https://www.hdfgroup.org/HDF5/>.
- Yang, Yimei, Jinping Liu, Yujun Yang, and Lihui Cen. 2024. “Dual-Stream Fusion Network with Multi-Head Self-Attention for Multi-Modal Fake News Detection.” *Applied Soft Computing* 167: 112358. <https://doi.org/https://doi.org/10.1016/j.asoc.2024.112358>.
- Zhang, Tian, Raghu Ramakrishnan, and Miron Livny. 1997. “BIRCH: A New Data Clustering Algorithm and Its Applications.” *Data Mining and Knowledge Discovery* 1 (2): 141–82. <https://doi.org/10.1023/A:1009783824328>.