

StackExchange Prediction and Analysis

ADM stackexchange project

Mateusz Mazur, Tomasz Makowski, Tomasz Kawiak

December 2, 2025

Table of contents



1 EDA

2 TAG Embeddings

3 TAG 2

4 Tag Prediction

EDA

- As of now, we have collected a dataset containing 100,000 questions from StackExchange (specifically StackOverflow), along with their associated metadata.

- As stated earlier, the dataset has the following features:

#	Column	Dtype
0	title	object
1	has_accepted_answer	bool
2	accepted_answer_score	float64
3	time_to_accepted_answer_hours	float64
4	question_score	int64
5	question_text	object
6	num_tags	int64
7	tags	object
8	accepted_answer_id	float64
9	accepted_answer_length_chars	float64
10	accepted_answer_length_tokens	float64

Deduplication

- Out of these questions, 8 were dropped for being duplicates.

Acceptance Analysis

- From those remaining, 39938 questions have an accepted answer, while 60054 do not.
- Interestingly, out of those accepted answers, only 12000 have `time_to_accepted_answer_hours` defined (i.e., non-null values).

Time to Accepted

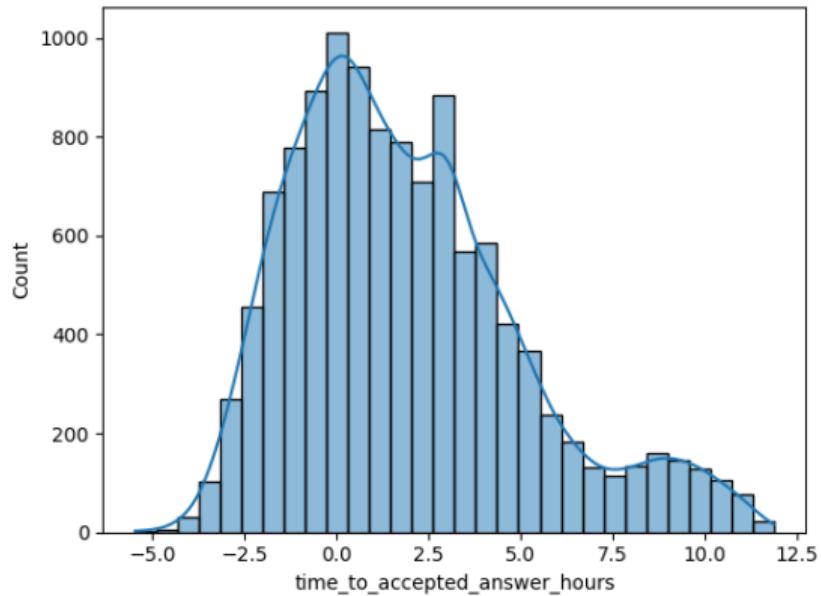


Figure 1: Histogram of log of time to accepted answer

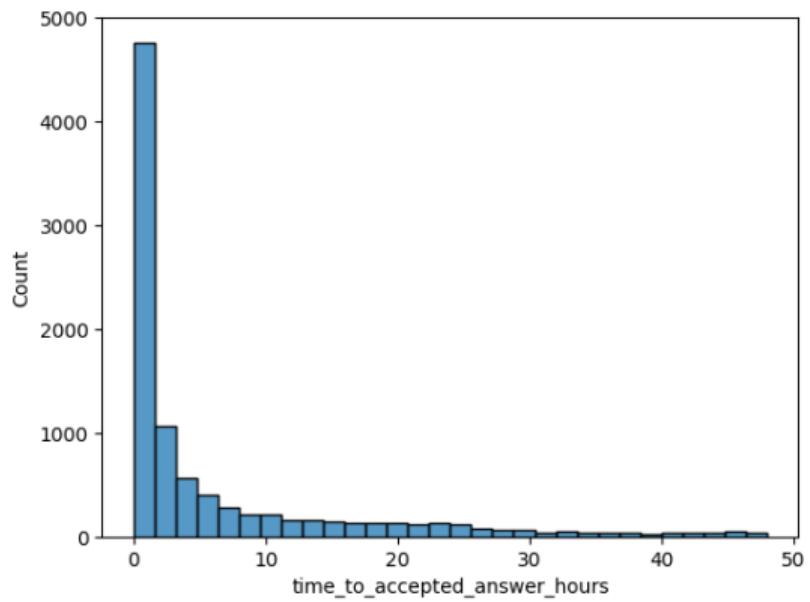


Figure 2: Histogram of time to accepted answer with time < 48

Tags

- The dataset contains a total of 7684 unique tags. The most common being:

Tag	Count
python	1528
c#	746
javascript	703
c++	689
java	592

- Since we are dealing with such a vast number of different tags, we propose the following approaches to investigate:
 - 1 Frequency-Based Filtering
 - ▶ Aside from the top N most common tags, one approach could limit the scope to a subset of tags relevant to our analysis – e.g. **top N programming languages**
 - 2 “Semantic Clustering”:
 - ▶ Use pre-trained embeddings to represent question descriptions in a vector space.
 - ▶ Apply clustering algorithms to group similar questions together based on their embeddings.

Tags EDA




Figure 3: Tag Frequency

We filtered the tags using the list of programming languages (from Wikipedia) and selected the most frequent ones:

```
['python', 'javascript', 'java', 'c#', 'php', 'html', 'c++', 'css', 'r']
```

Key observations about the resulting subset:

- **Subset size:** 42,037 questions
- **Class balance:** The distribution across these languages is relatively balanced (see last plot), and can be further balanced if needed.
- **Multi-label cases:** Some questions are tagged with multiple programming languages.

Distribution of programming language tags per question:

# Tags	# Questions
1	38,547
2	3,015
3	464
4	11

- Finally, we analyzed the distribution of the scores of the questions in the dataset (already of the selected programming-language-tagged subset).

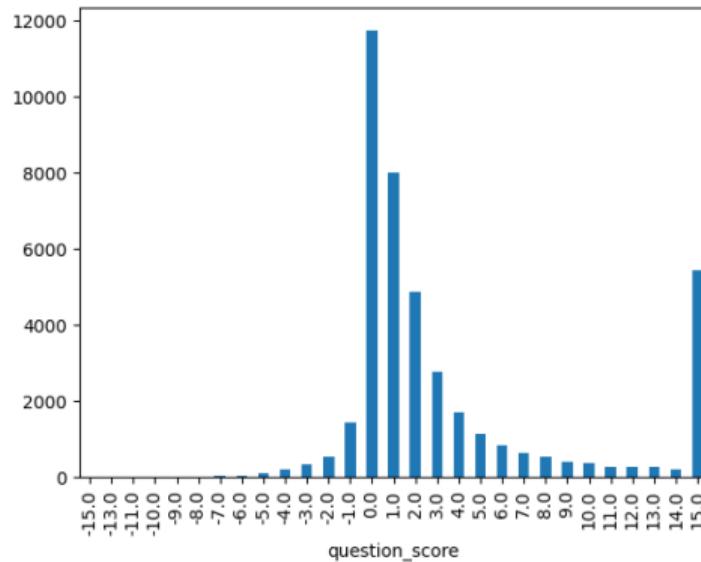


Figure 4: Question Score Histogram

Based on the histogram, we proposed the following score classes:

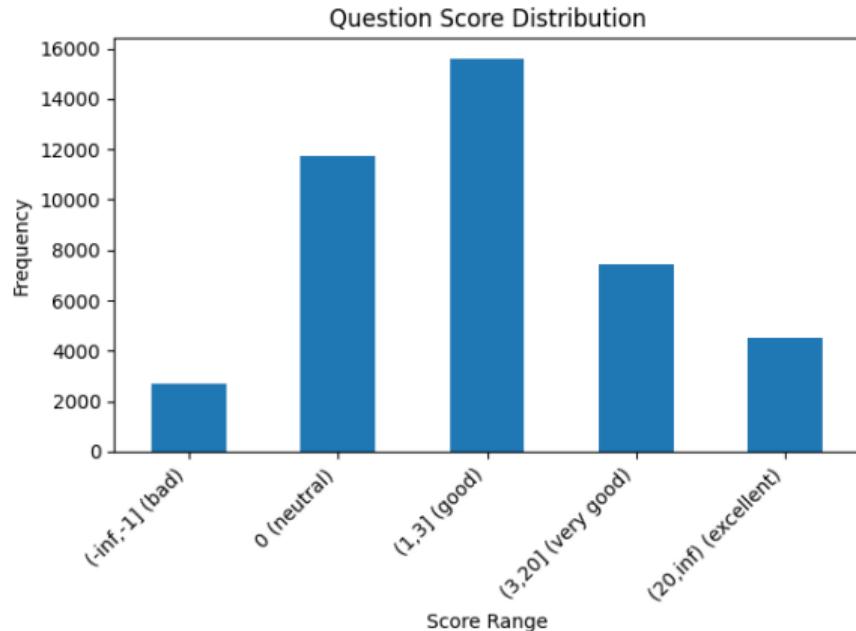


Figure 5: Question Score Distribution

Proportion of Questions where Tag(s) Appears in Text

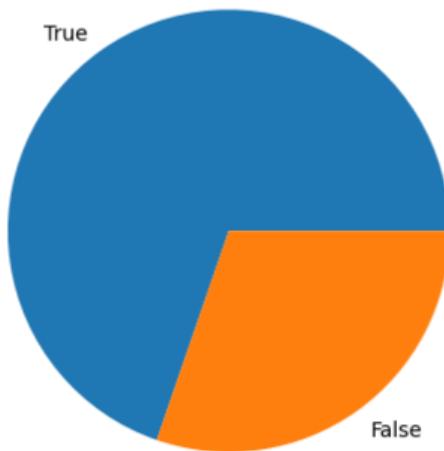


Figure 6: Proportion of Questions where Tag(s) appears in text

- the data about the accepted answer is provided for only small fraction of the data.
- there are too many tags to perform the classification using all of them
- however, we can choose a certain subset (e.g. most popular programming languages) and focus on it
- we can also try to predict the class of the score of the answer (integer score mapped to classes of uneven frequencies)

TAG Embeddings

Initial problem



- Having 22753 unique tags poses a significant problem for our dataset,
- each questions can have multiple tags (for example python,pandas,...)

Embeddings

- We've decided to embed tags, as well as question texts using qwen3-embedding:8b model, translating each tag into a 4096 vector.
- This resulted in a dataframe of shape 22753x4096 vector, which introduced more issues:
 - ▶ Lack of memory to process this data (the dtype is f64)
 - ▶ Too large dimensions




Figure 7: Question embeddings cosine similarity.

Dimensionality reduction



- Initial approach considered HDBSCAN (Hierarchical Density-Based Spatial Clustering of Applications with Noise), but due to memory required to perform this operation (over 64GB of RAM)
- Later on we've moved to a Birch algorithm, with HDBSCAN as well as AgglomerativeClustering which got unsatisfactory results

- After reading forums, we've decided to try UMAP (Uniform Manifold Approximation and Projection) which sounded ideal for our problem, since it can be used for general non-linear dimension reduction.

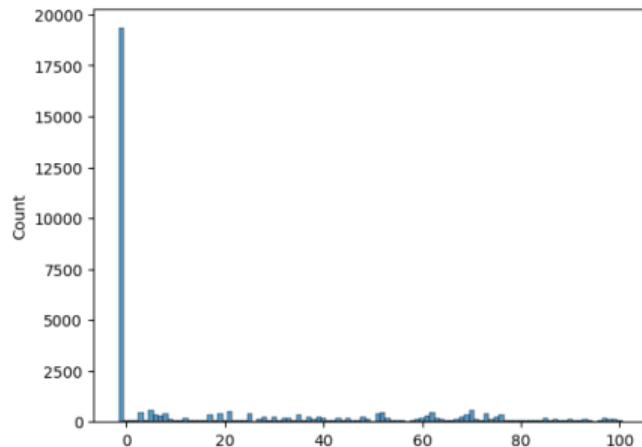


Figure 8: Recurring problem during UMAP and other clustering algorithms




Figure 9: Naiive UMAP reducing tag embeddings to 2 components.




Figure 10: Histogram of labels resulting from UMAP reduction coupled with HDBSCAN

- A hyperparameter optimization framework which we hoped would find optimal parameters for our dimensionality reduction task.
- We need a score which we can maximize/minimize during dimension reduction:
 - ▶ Caliński-Harabasz index
 - ratio of the between-cluster separation to the within-cluster dispersion, normalized by their number of degrees of freedom.
 - ▶ Silhouette score
 - ▶ Cluster persistence
 - stability of each cluster, indicating well-defined grouping

UMAP Optimization



- UMAP coupled with HDBSCAN
- Optimized hyperparameters:
 - ▶ `umap_n_components`
 - ▶ `umap_n_neighbors`
 - ▶ `umap_min_dist`
 - ▶ `umap_metric`
 - ▶ `hdb_min_cluster_size`
 - ▶ `hdb_min_samples`
 - ▶ `hdb_cluster_selection_method`
 - ▶ `hdb_cluster_selection_epsilon`

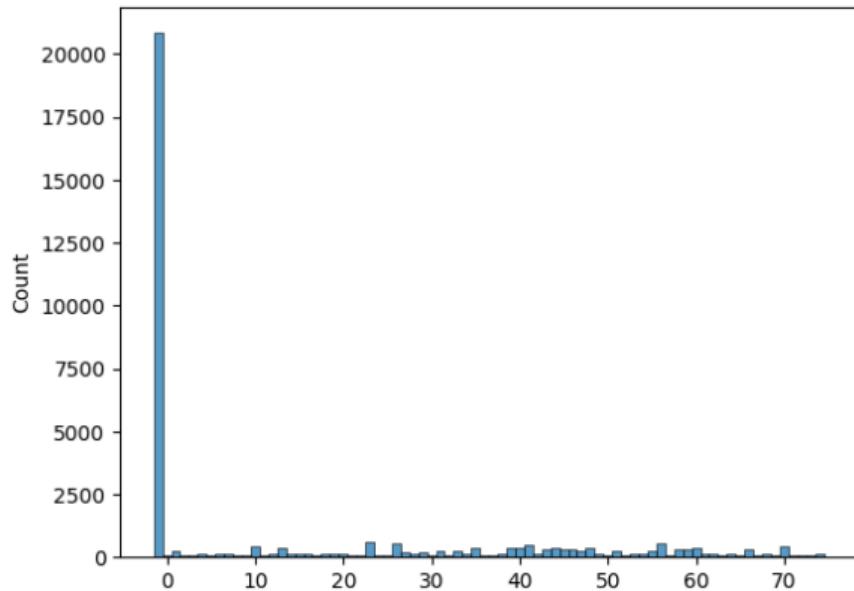


Figure 11: Results of optuna training

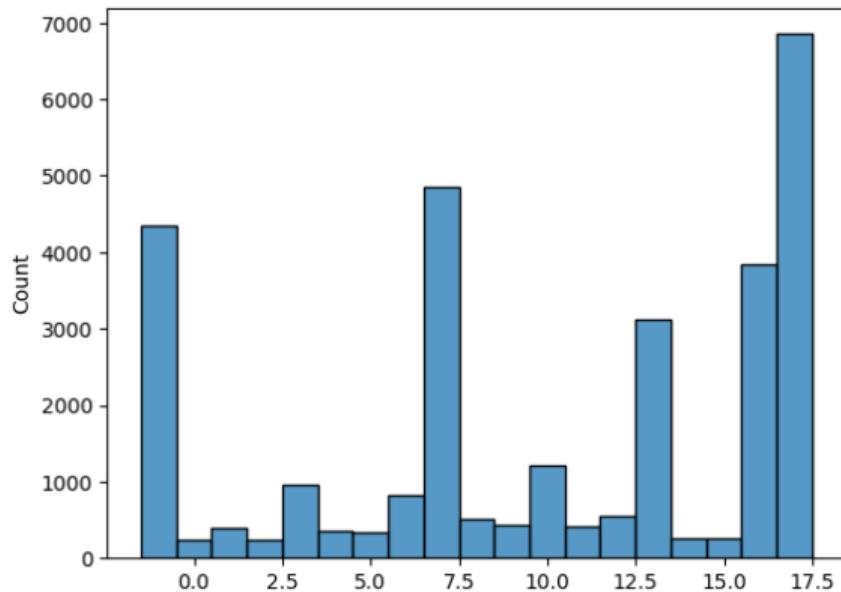


Figure 12: Results of Optuna optimization of only HDBSCAN hyperparameters search ran on UMAP reduction to 5 components

- We have to find better approach
- Possible culprits:
 - ▶ For UMAP to work as intended 'The data must be uniformly distributed on Riemannian manifold'
 - ▶ additionally 'The Riemannian metric is locally constant (or can be approximated as such)'
 - embeddings generally create spaces with varying local geometry depending on the semantic density of the training data.
 - ▶ Too few optimization runs.

Recursive Embedding and Clustering

- Based on Spotify's blog.

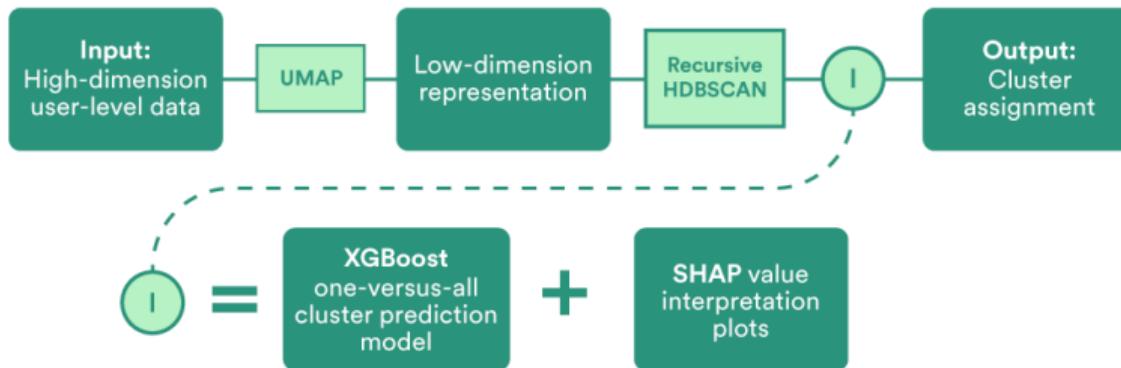


Figure 13: Diagram representing the complete recursive embedding and clustering process. Source

Recursive KMEANS clustering

- Since properly implementing HDBSCAN in the context of our problem is still a standing question, we've decided to use KMEANS.
- Initially, we've specified the clustering hierarchy structure 100,100.
- After finding those centroids, we've taken their mean embeddings and selected 'most descriptive' tags for each cluster (based on cosine/euclidean metric).
- euclidean metric worked 'better' in the case of KMEANS.

		cluster_mean	top_5_popular_tags
	facebook	[0.0052028694, -0.002167758, -0.017349796, -0....	[instagram-api, instagram, instagram-graph-api...
	pyobject	[0.004231446, 0.0026400657, -0.0142704435, -0....	[python, pandas, jupyter-notebook, jupyter, py...
google-workspace-add-ons		[0.022996485, 0.0034478805, -0.014011028, -0.0...	[google-merchant-center, google-signin, google...
.net		[0.008871326, 0.00042595304, -0.003047627, -0....	[c#, wpf, .net, blazor, asp.net-core]
coverflow		[0.018253814, 0.0024098884, 0.0059977337, -0.0...	[background-color, css, tailwind-css, css-posi...
azure		[0.02082736, 0.0026177948, 0.0014911512, -0.00...	[azure-pipelines, azure, azureservicebus, azur...
importerror		[0.017333947, 0.00064369285, 0.0025193274, -0....	[segmentation-fault, try-catch, connection-ref...

Figure 14: Initial results

Recursive Spherical KMEANS clustering



- Since spherical KMEANS is better suited for cosine similarity and for text data in general, we've decided to try it as well. Based on a Accelerating Spherical k-Means paper [2].

	cluster_mean	top_5_popular_tags
facebook	[0.004879954, -0.0019050128, -0.016777342, -0....	[instagram-api, instagram, instagram-graph-api...
integer	[0.027635492, 0.0065359636, -0.0029253308, -0....	[intervals, differential-equations, math, maxi...
storage	[0.008282, -0.016636355, -0.0092419945, -0.004...	[dicomweb, remotestorage, storage, multipartfo...
tf-idf	[0.027263727, -0.007913517, 0.0023698097, -0.0...	[wpml, locale, dictionary, spacy, microsoft-tr...
qt	[0.01614129, 0.005696066, 0.016760174, -0.0084...	[qt, qml, pyqt, qasync, qt6]

Figure 15: Sample results

Results

- To verify the quality of new clusters, we've used a very basic predictor based on the dot product of normalized question embeddings and tag cluster embeddings.
 - ▶ We've taken 2 most similar clusters for each question.




Figure 16: Results of top 2 results for .net

	tag1	tag2
count	780	780
unique	780	780
top	(mongodb, 0.5270306638492696)	(hyperfilesql, 0.4293950990405074)
freq	1	1

Figure 17: Most popular results for mongodb

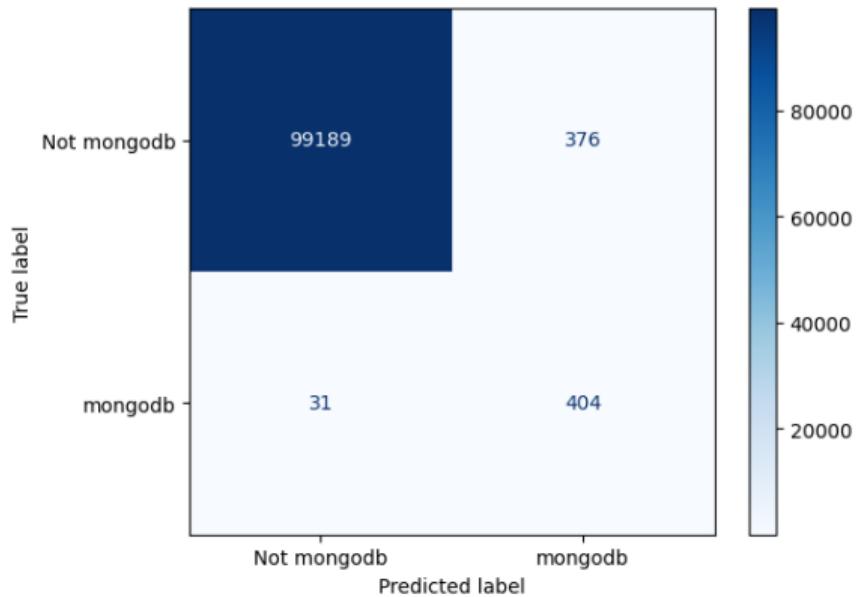


Figure 18: Results of top 2 results for mongoDB

TAG 2

- Problem: 22,753 unique tags is too large for direct multi-label classification.
- Solution: Reduce tags to 100 clusters using *Recursive Spherical K-Means* (with cosine similarity), then classify into these clusters.

Reduction & Baseline

- Our dataset has 100k records many of which have multiple tags.
- The idea is that we can remove tags with low frequency and remove records that become tag-less after this operation.
- Aforementioned reduction for the threshold of 100 resulted in 90,205 records and 411 unique tags.

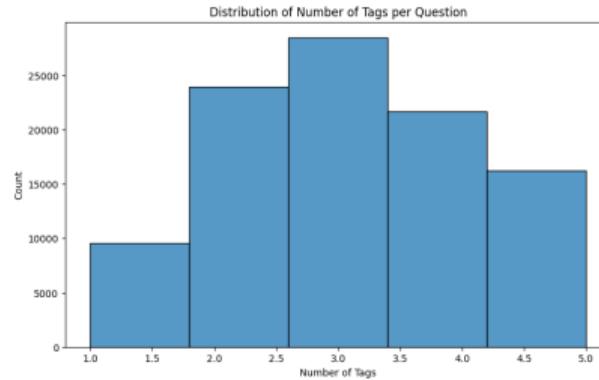


Figure 19: Distribution of number of tags per question

- With the resulting tags, we performed Recursive Spherical K-Means to reduce the 411 tags to 100 centroids.
- The outlier with 15 tags corresponds to a cluster with following tags: database, mongodb, sql, elasticsearch, firebase, postgresql, sqlite, mysql, jdbc, supabase, sql-update, redis, hibernate, sqlalchemy, snowflake-cloud-data-platform

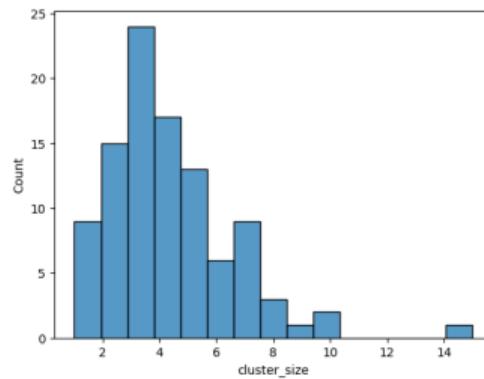


Figure 20: Distribution of number of clusters per centroid

- For the 9787 orphan tags (those with insufficient frequency), we assigned them to the nearest centroid based on cosine similarity of their embeddings.
- For example:

```
closest_centroid_tag(tags,centroid_tags, 'ntp')
('datetime', np.float64(0.8098148848579392))
```

- This resulted in the final set of 100 centroids covering all 22,753 original tags, without losing any data.

Why Neural Networks?



- 1 **Multi-Label Classification:** NNs naturally handle multi-label outputs (e.g., using Sigmoid activations + BCE Loss).
- 2 **Feature Interaction:** Deep learning models can learn complex non-linear interactions between Title and Body embeddings.
- 3 **Custom Loss Functions:** Easier to implement losses like Asymmetric Loss to handle class imbalance.

Tag Prediction

- XGBoost stands for “Extreme Gradient Boosting”, where the term “Gradient Boosting” originates from the paper *Greedy Function Approximation: A Gradient Boosting Machine*, by Friedman.
- It is an ensemble learning method that builds multiple weak learners (decision trees) sequentially, where each tree attempts to correct the errors of the previous ones.
- Task: **predict the question's centroid tag based on question text embedding**.

- Our questions still have multiple tags, which drastically complicates the problem.
 - ▶ To solve this, we first took the original tags of each question, mapped them to their respective centroids, and selected the “best” centroid (the one with the most original tags assigned to it) as the target label.

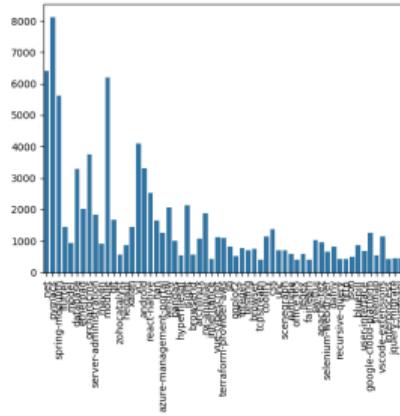


Figure 21: Resulting count of questions per centroid

- To combat the unbalanced distribution of centroid tags, we applied class weights inversely proportional to their frequencies during training, using `sklearn.utils.class_weight.compute_class_weight` and passing these weights to XGBoost's `scale_pos_weight` parameter.
- Learning process took 599 minutes.
- Accuracy: 0.6928123848138592
- F1 Weighted: 0.6882358640179885

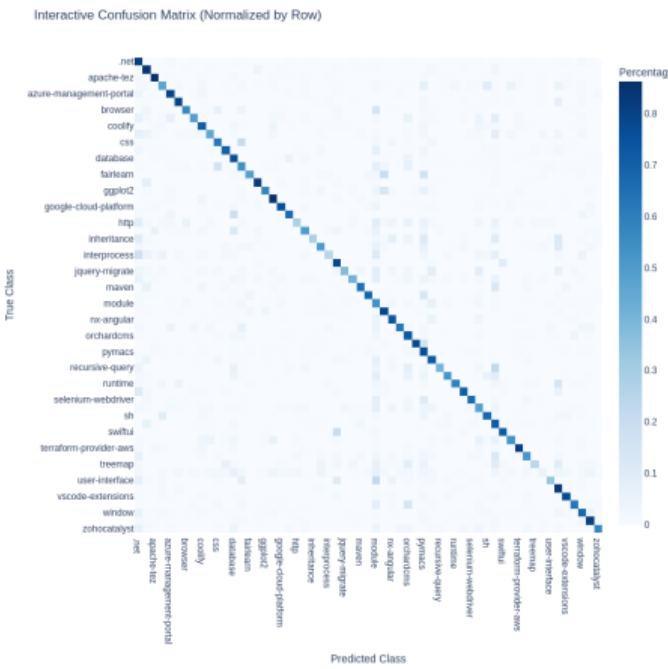


Figure 22: Confusion Matrix of XGBoost Classifier

Limitations of XGBoost Approach

- As a tree-based model, XGBoost treats input dimensions largely independently.
- It struggles to capture the complex, non-linear semantic interactions present in text embeddings.
- Training on 4096-dimensional dense embeddings is computationally expensive for gradient boosting trees.
- It was impossible to tune hyperparameters effectively due to long training time.

- For better modeling of the multi-label nature of the problem, we turned to Neural Networks using PyTorch.
- We can easily design architectures for multi-label classification and implement custom loss functions to handle class imbalance.
- We can leverage **Attention Mechanisms** and other fancy architectures for our predictive model.

Baseline: Simple MLP



■ Architecture:

```
BaselineNetwork(  
    (fc1): Linear(in_features=4096, out_features=2048)  
    (relu): ReLU()  
    (dropout): Dropout(p=0.3)  
    (fc2): Linear(in_features=2048, out_features=1024)  
    (dropout2): Dropout(p=0.3)  
    (fc3): Linear(in_features=1024, out_features=100)  
)
```

■ Input: question_text_embedding

■ Loss: BCEWithLogitsLoss

- ▶ Since question can have a very popular, and very rare tag (e.g. .net, nvim), we can't just weigh all tags equally. Instead, we will penalize by BCEWithLogitsLoss for missing a rare tag much more heavily than for missing a popular tag.

- Training:
 - ▶ 50 epochs ~6 minutes
 - ▶ F1 (weighted): $\approx 67.9\%$
- Properly models multi-label outputs (unlike a multiclass XGBoost mapping)
- Low compute and quick iteration for ablations
- Slightly underperforms XGBoost baseline on this dataset, but performs multi-label classification instead of basic classification.



Figure 23: Per-Tag Performance of Baseline MLP

Per-Tag Performance: Python

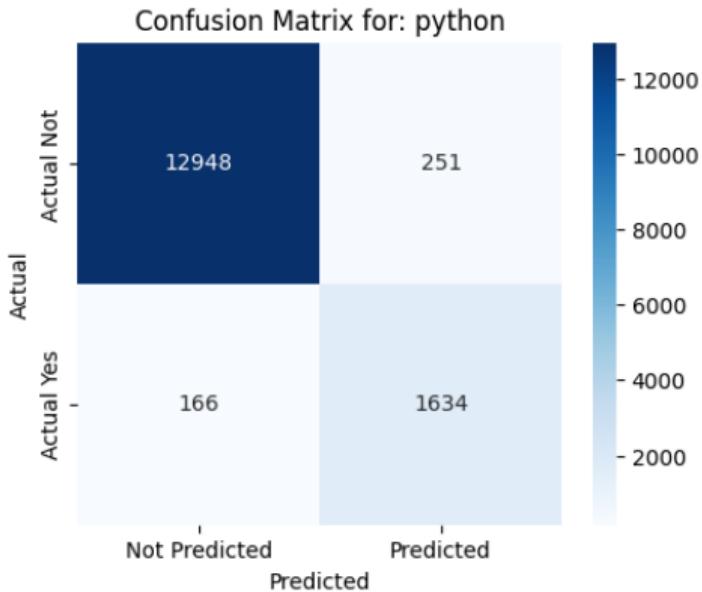


Figure 24: Confusion Matrix for Python only

- Inspired by ‘Dual-stream fusion network with multi-head self-attention for multi-modal fake news detection’ [3].
- Instead of processing text and images as in the paper above, we will map two streams:
 - ▶ question_text embeddings,
 - ▶ title embeddings,
- and fuse them using multi-head self-attention to capture correlation between the two modalities.
- As a fusion, we will use the paper’s proposed MHSA Fusion strategy to let the Title “attend” to the question body and vice versa.

Dual Stream Fusion Network (DSF)



- Process Title and Body separately and fuse them.
- **Architecture:**
 - ▶ **Stream 1:** Title Embedding -> Projection -> LayerNorm
 - ▶ **Stream 2:** Body Embedding -> Projection -> LayerNorm
 - ▶ **Fusion:** Multi-Head Self-Attention (MHSA) on stacked features.
- Since the shape of both embeddings is the same, we have the same architecture for both streams:

```
(title_proj): Sequential(  
    (0): Linear(in_features=4096, out_features=1024, bias=True)  
    (1): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)  
    (2): GELU(approximate='none')  
    (3): Dropout(p=0.1, inplace=False)
```

- The initial results were disappointing- model overfitted very quickly and proved immediately useless on validation set.
- To mitigate overfitting, we:
 - ▶ Increased dropout rate from 0.1 to 0.5.
 - ▶ Changed optimizer from Adam to AdamW.
 - AdamW decouples weight decay from the gradient update, leading to better generalization.
 - ▶ Implemented **Asymmetric Loss** [1] to handle label imbalance.

Asymmetric Loss (ASL)

- In multi-label classification with 100 classes, for any given question, most labels are **Negative**. The model can get high accuracy by just predicting “0” for everything.
- To counter this, ASL:
 - ▶ Separates the treatment of positive and negative samples.
 - Applies different focusing parameters to down-weight easy negatives more aggressively.
 - Keeps the weight high for positive samples to ensure they are learned well.
 - ▶ Focuses training on hard negatives and positive samples.
 - ▶ Essentially, mistakes on positives matter more than mistakes on easy negatives (based on the probability score output by the model).
- Parameters: $\gamma_- = 4$, $\gamma_+ = 1$.

- Aside from the loss function, we've added:
 - ▶ ReduceLROnPlateau scheduler to reduce learning rate on validation F1 plateau.
 - ▶ As earlier mentioned, AdamW optimizer with weight decay of 1e-2.
- Training for 50 epochs took ~20 minutes.
- F1 score (micro) on validation: 0.713
- F1 score (weighted) on validation: 0.711

Hedgehog Anatomy




Figure 25: Per-Tag Performance & Multilabel Confusion Matrices of DSF_MHSA_Classifier

Dual Stream Fusion with Cross-Attention



- Some sort of overfitting still occurred, so we tried replacing the MHSA fusion with Cross-Attention:
 - ▶ Instead of both streams attending to each other equally, we let the Title embedding "query" the Body embedding.
- Intuition:
 - ▶ *"Given this Title, which parts of the Body Embedding are relevant?"*
 - ▶ Title is short and dense with information.
 - ▶ Body is longer and more verbose.
 - ▶ Letting Title attend to Body helps focus on relevant parts of the Body text.
- Same as before, we used ASL loss to combat label imbalance. This time with $\gamma_- = 2$, $\gamma_+ = 1$.
- To further reduce overfitting, we also applied Manifold Mixup on the embeddings:
 - ▶ It creates new training samples by interpolating between the embeddings (and tags) of random pairs.
 - ▶ Enforces smoother decision boundaries in the latent space, improving generalization.

- Training for 100 epochs took ~60 minutes.
- Used OneCycleLR scheduler.
- Implemented basic Curriculum Learning by starting with smaller dropout (0.25)
- F1 score (micro) on validation: 0.725291274763489
- F1 score (weighted) on validation: 0.7196345470456641



Figure 26: Per-Tag Performance & Multilabel Confusion Matrices of DSF_CrossAttn_Classifier

Summary of Results



Model	F1 (weighted)
XGBoost	0.6882
Baseline MLP	0.6790
DSF with MHSA Fusion	0.7110
DSF with Cross-Attention Fusion	0.7196

Thank you for your attention

Questions

References

References

- [1] Ben-Baruch, E., Ridnik, T., Zamir, N., Noy, A., Friedman, I., Protter, M. and Zelnik-Manor, L. 2021. Asymmetric loss for multi-label classification.
- [2] Schubert, E., Lang, A. and Feher, G. 2021. Accelerating spherical k-means. *Similarity search and applications*. Springer International Publishing. 217–231.
- [3] Yang, Y., Liu, J., Yang, Y. and Cen, L. 2024. Dual-stream fusion network with multi-head self-attention for multi-modal fake news detection. *Applied Soft Computing*. 167, (2024), 112358. DOI:<https://doi.org/10.1016/j.asoc.2024.112358>.