

# StackExchange Prediction and Analysis

## ADM stackexchange project

Mateusz Mazur, Tomasz Makowski, Tomasz Kawiak

Faculty of Electrical Engineering, Automation, Computer Science and Biomedical Engineering, AGH  
Field of study: Computer Science and Intelligent Systems  
Specialization: Artificial Intelligence and Data Analysis

November 6, 2025

# Table of contents



**1** EDA

**2** References

# EDA

# Overview



- As of now, we have collected a dataset containing 100,000 questions from StackExchange (specifically StackOverflow), along with their associated metadata.

- As stated earlier, the dataset has the following features:

#	Column	Dtype
0	title	object
1	has_accepted_answer	bool
2	accepted_answer_score	float64
3	time_to_accepted_answer_hours	float64
4	question_score	int64
5	question_text	object
6	num_tags	int64
7	tags	object
8	accepted_answer_id	float64
9	accepted_answer_length_chars	float64
10	accepted_answer_length_tokens	float64

## Deduplication

- Out of these questions, 8 were dropped for being duplicates.

## Acceptance Analysis

- From those remaining, 39938 questions have an accepted answer, while 60054 do not.
- Interestingly, out of those accepted answers, only 12000 have `time_to_accepted_answer_hours` defined (i.e., non-null values).

# Time to Accepted

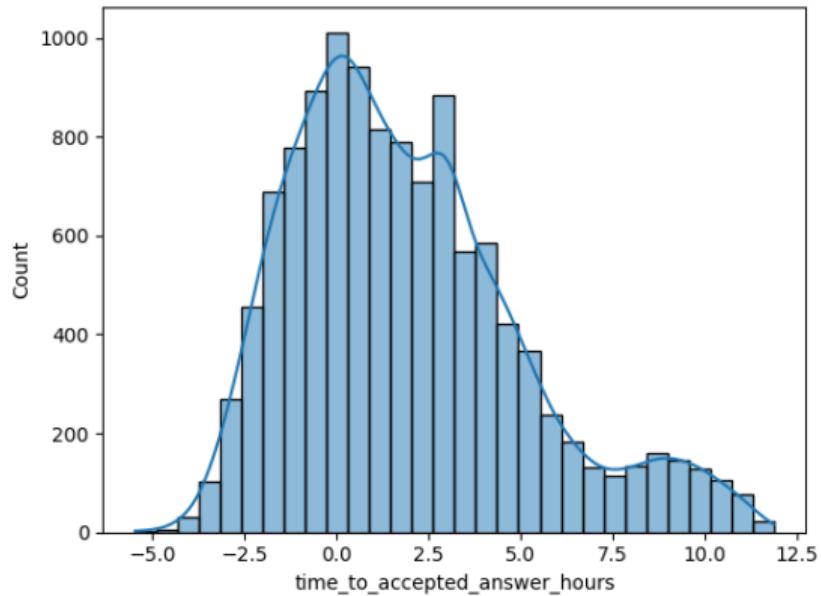


Figure 1: Histogram of log of time to accepted answer

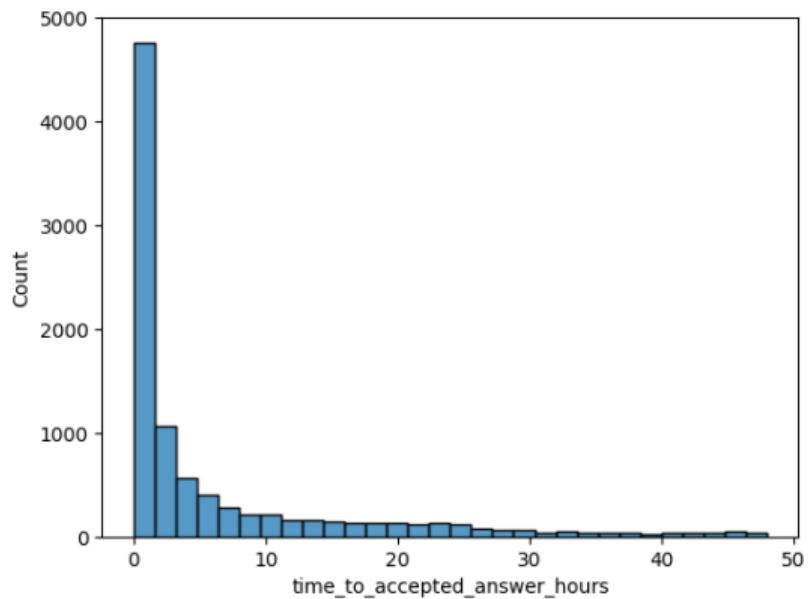


Figure 2: Histogram of time to accepted answer with time < 48

# Tags

- The dataset contains a total of 7684 unique tags. The most common being:

Tag	Count
python	1528
c#	746
javascript	703
c++	689
java	592



- Since we are dealing with such a vast number of different tags, we propose the following approaches to investigate:
  - 1 Frequency-Based Filtering
    - ▶ Aside from the top  $N$  most common tags, one approach could limit the scope to a subset of tags relevant to our analysis – e.g. **top  $N$  programming languages**
  - 2 “Semantic Clustering”:
    - ▶ Use pre-trained embeddings to represent question descriptions in a vector space.
    - ▶ Apply clustering algorithms to group similar questions together based on their embeddings.

# Tags EDA

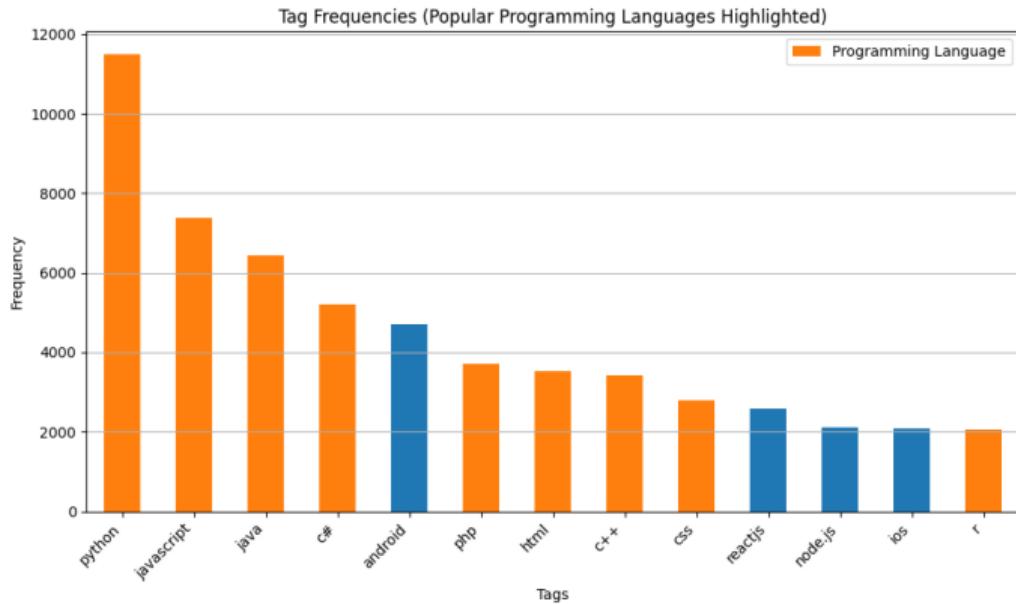


Figure 3: Tag Frequency

We filtered the tags using the list of programming languages (from Wikipedia) and selected the most frequent ones:

```
['python', 'javascript', 'java', 'c#', 'php', 'html', 'c++', 'css', 'r']
```

### Key observations about the resulting subset:

- **Subset size:** 42,037 questions
- **Class balance:** The distribution across these languages is relatively balanced (see last plot), and can be further balanced if needed.
- **Multi-label cases:** Some questions are tagged with multiple programming languages.

## Distribution of programming language tags per question:

# Tags	# Questions
1	38,547
2	3,015
3	464
4	11

- Finally, we analyzed the distribution of the scores of the questions in the dataset (already of the selected programming-language-tagged subset).

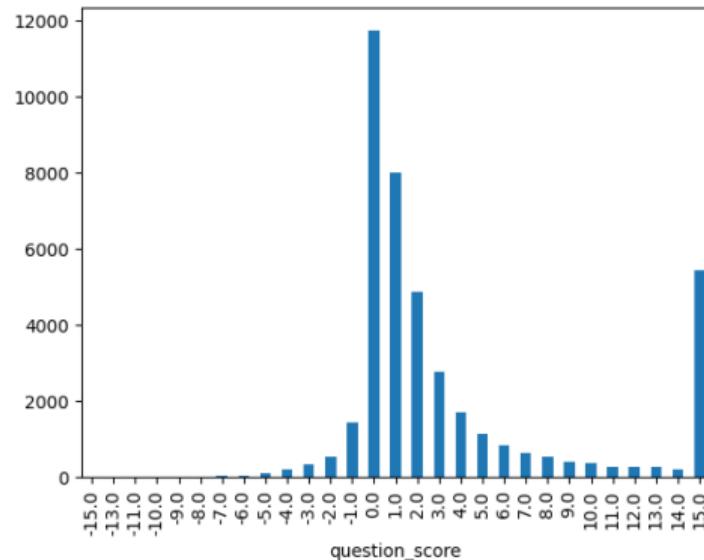


Figure 4: Question Score Histogram

Based on the histogram, we proposed the following score classes:

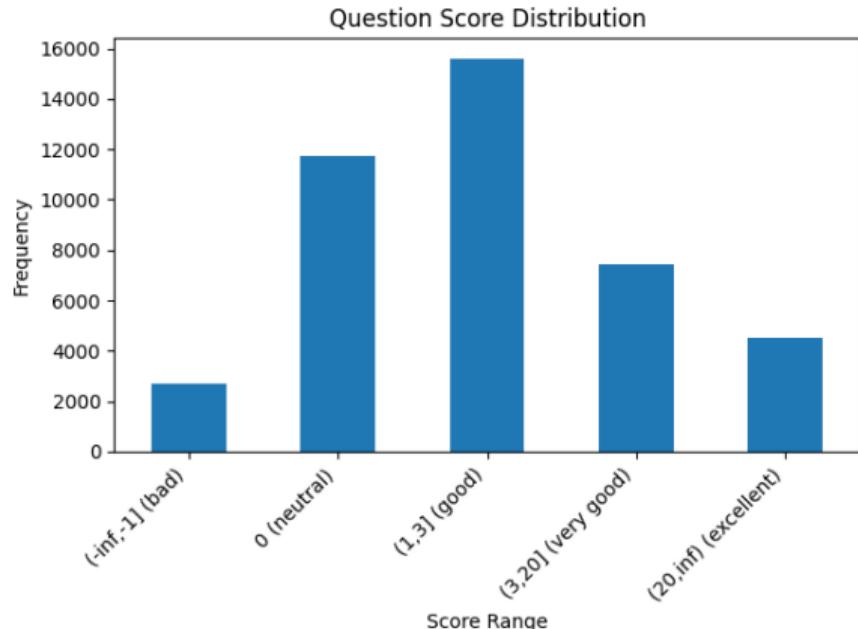


Figure 5: Question Score Distribution

Proportion of Questions where Tag(s) Appears in Text

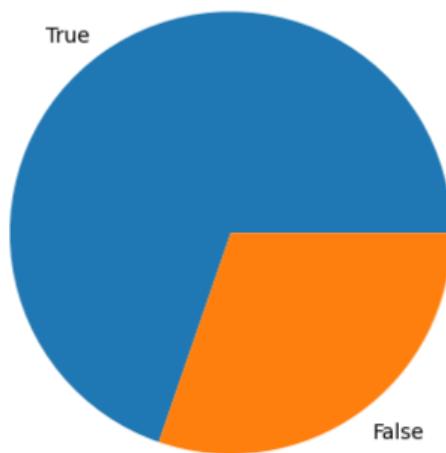


Figure 6: Proportion of Questions where Tag(s) appears in text

# Conclusions

- the data about the accepted answer is provided for only small fraction of the data.
- there are too many tags to perform the classification using all of them
- however, we can choose a certain subset (e.g. most popular programming languages) and focus on it
- we can also try to predict the class of the score of the answer (integer score mapped to classes of uneven frequencies)

Thank you for your attention

## Questions

## References